

Sampler Proliferation

Sampler Proliferation

- Many samplers collect the same or nearly the same information
 - procstat, procstat2
 - app_sampler, linux_proc_sampler
 - slurm_sampler, slurm_sampler2
- Creates user confusion: which variant should I use and why
- Lots of duplicated code
- Doubles or worse the software maintenance costs
 - If I want to add a new metric to sampler A, must I add it to B too?
 - We have to test both samplers
 - We have to document both samplers
- Many samplers are no longer used

Possible solutions

- If the schema is largely identical, add the new feature to the existing sampler and select behaviors with configuration options
- Pull common/duplicated code out as appropriate and bind the common code into both samplers in the Makefile
- Determine which samplers are no longer used and remove them from the source tree

Deprecating Samplers

- Process
- Time frame
- Selection Criteria
- Feedback from users

Q&A